

E6156 – Topics in SW Engineering: Microservices and Cloud Applications

Lecture 6: Web API, Pub/Sub, Messaging

Donald F. Ferguson, dff@cs.columbia.edu

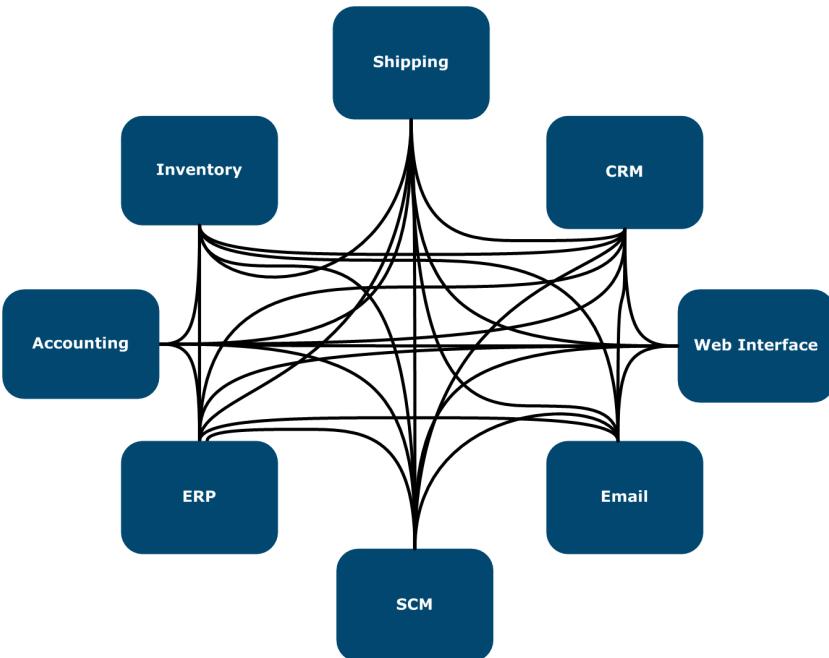
Introduction

Contents

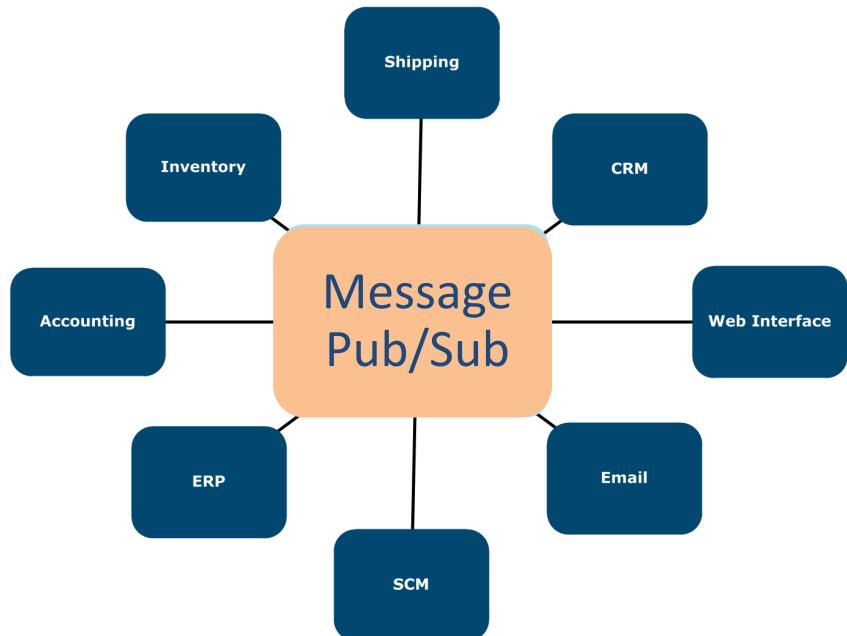
- Whiteboard discussion of:
 - Base project topics to date.
 - What we should try to accomplish in foundation to enable the simple application you are building.
- Web APIs:
 - Profile and address verification example:
 - End-to-end flow.
 - Client side
 - Stripe: Credit card payment overview (if you want to simulate it).
- Asynchronous microservice assembly:
 - Almost all programming you have done in university is tightly coupled *call-return*.
 - An extremely, if not most common approach, to coupling systems is loosely-coupled:
 - Publish/Subscribe
 - Message driven queuing/message driven processing
 - We will get a feel for pub/sub by including in email verification.

Integration/Assembly Patterns

Classic Approach



Pub/Sub or Message Hub

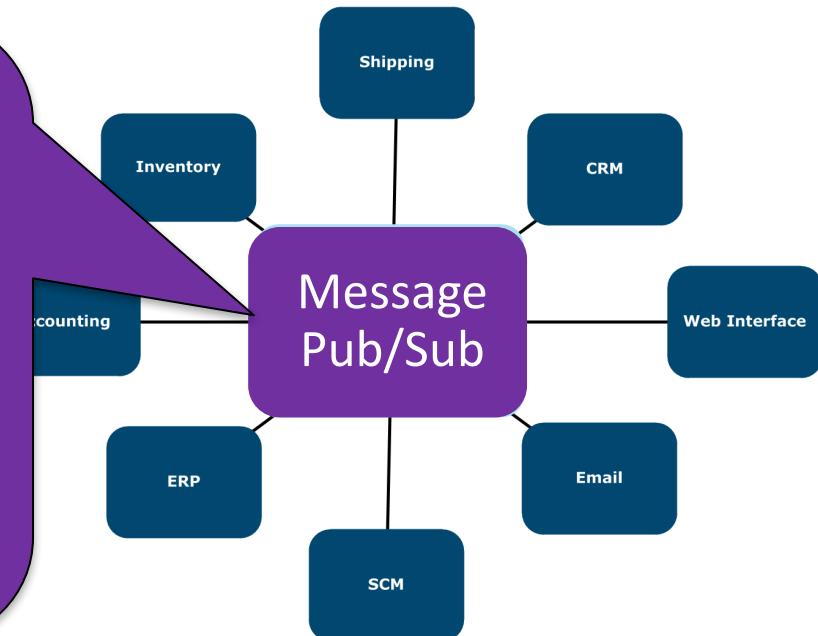


Integration/Assembly Patterns

Classic Approach

- The logical model is a “hub” supporting
 - Topics
 - Queues
- The physical model is a robust, multi-system, highly available and high performance infrastructure

Pub/Sub or Message Hub



Web APIs

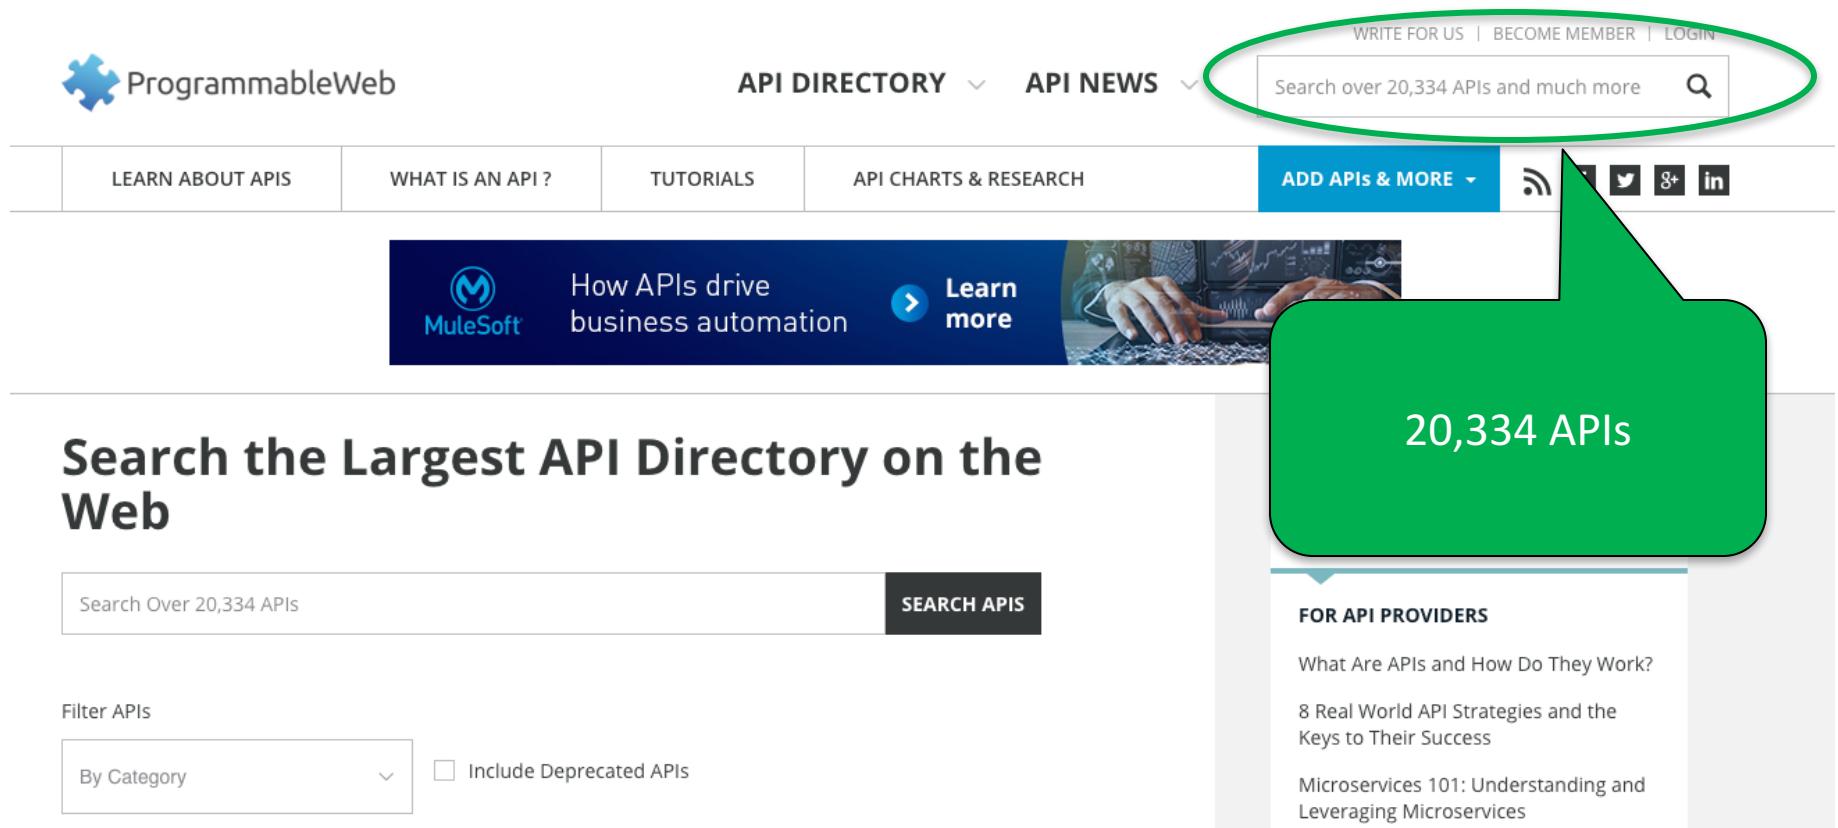
API Economy

Concept

Extensive Use of Web APIs



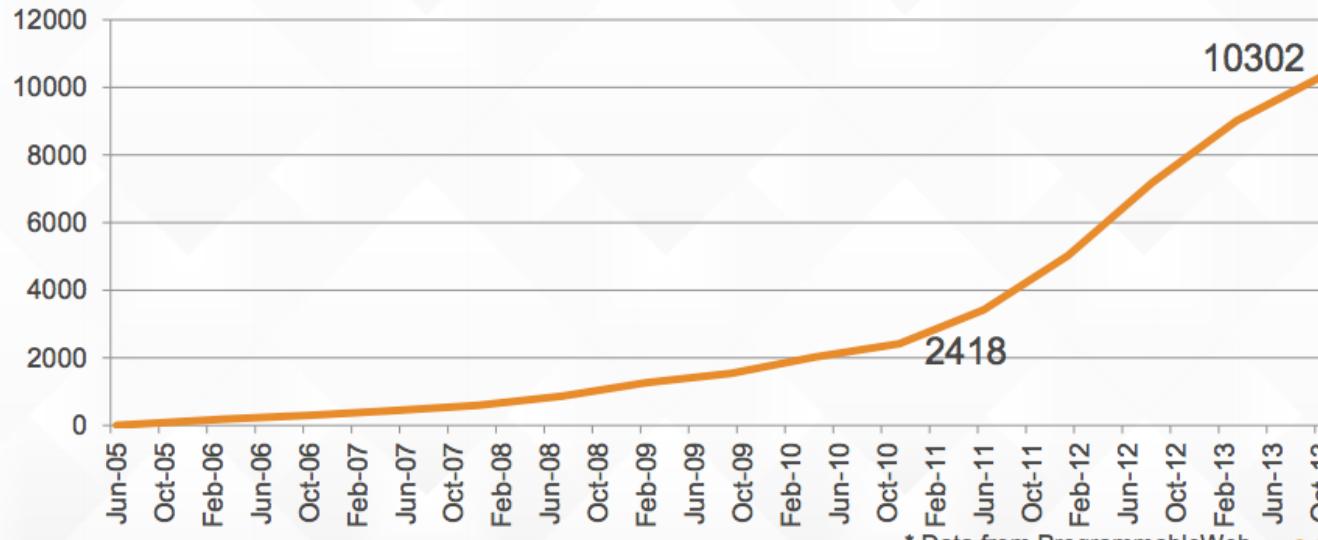
Programmable Web – An Example



The screenshot shows the ProgrammableWeb website. At the top, there is a navigation bar with links for "API DIRECTORY", "API NEWS", "WRITE FOR US", "BECOME MEMBER", and "LOGIN". A search bar is prominently displayed, with a green oval highlighting it. Below the navigation bar, there are links for "LEARN ABOUT APIs", "WHAT IS AN API?", "TUTORIALS", and "API CHARTS & RESEARCH". A "ADD APIs & MORE" button is also present. On the right side of the header, there are social media icons for RSS, Twitter, Google+, and LinkedIn. A banner for MuleSoft is visible, featuring the text "How APIs drive business automation" and a "Learn more" button. A green callout bubble on the right side of the page contains the text "20,334 APIs". The main content area features a large heading "Search the Largest API Directory on the Web" and a search bar with the placeholder "Search Over 20,334 APIs". Below the search bar are filter options for "By Category" and "Include Deprecated APIs". To the right, there is a section titled "FOR API PROVIDERS" with links to articles: "What Are APIs and How Do They Work?", "8 Real World API Strategies and the Keys to Their Success", and "Microservices 101: Understanding and Leveraging Microservices".

API proliferation

The number of published APIs is growing rapidly



Goals of API Design

Nevertheless, many API programs are growing out of previous SOA initiatives. Web services focused on internal or partner integrations are being opened up to developers—both within and outside the enterprise. During this process, it is important for API designers to remember that an API program has drivers and requirements quite different from the ones that initially led enterprises to open their IT assets via Web services.

With this in mind, the broad goals of API design in general can be defined as:

- Enabling self-service for app developers and app users alike
- Reducing barriers to accessing valuable enterprise resources
- Prioritizing the needs and preferences of client app developers
- Encouraging collaboration between and among internal and external resources
- Addressing the security and scaling issues of exposing IT assets to the open market

Above all, API design must be focused on maximizing the business value of the interface. In part two, we will take a closer look at how APIs add value to the business.

The API economy is an enabler for turning a business or organization into a platform.

We live in an API economy, a set of business models and channels based on secure access of functionality and exchange of data. APIs make it easier to integrate and connect people, places, systems, data, things and algorithms, create new user experiences, share data and information, authenticate people and things, enable transactions

and algorithms, leverage third-party algorithms, and create new product/services and business models.

“The API economy is an enabler for turning a business or organization into a platform.” said **Kristin R. Moyer**, vice president and distinguished analyst at Gartner. “Platforms multiply value creation because they enable business ecosystems inside and outside of the enterprise to consummate matches among users and facilitate the creation and/or exchange of goods, services and social currency so that all participants are able to capture value.”

Uber, for instance, is an example of a business built on a platform because it leverages Google Maps through an API to enable its entire business model of matching drivers who have a vehicle with passengers who need a ride. Walgreens offers an API for its in-store photo printing services that enables others to offer photo apps on its platform. It moves from being a photo printer to being a photo platform.

2017 Is Quickly Becoming The Year Of The API Economy



Louis Columbus, CONTRIBUTOR
[FULL BIO](#) ▾

Opinions expressed by Forbes Contributors are their own.

<https://www.forbes.com/sites/louiscolumbus/2017/01/29/2017-is-quickly-becoming-the-year-of-the-api-economy/#4b3806fe6a41>

This year more CIOs will have their bonuses tied to how many new business models they help create with existing and planned IT platforms than ever before. This trend will accelerate over the next three years. CIOs and IT staffs need to start thinking about how they can become business strategists first, technicians and enablers of IT second. CIOs must create and launch new business models faster to keep their companies competitive. APIs are the fuel helping to make this happen.

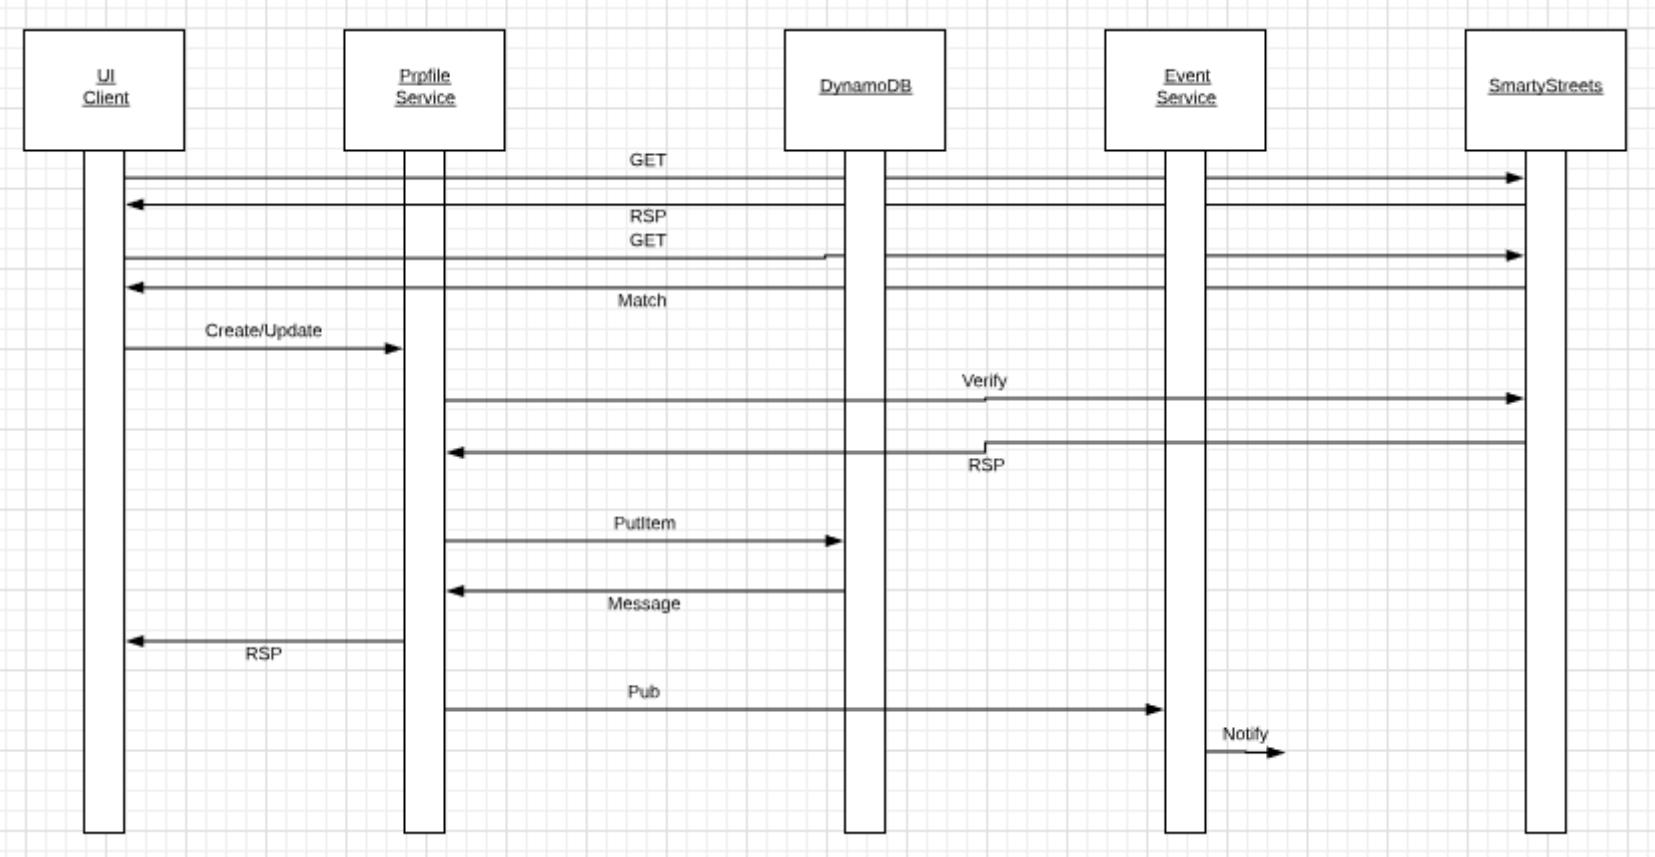
The Urgency To Create New Business Models Is Driving API Proliferation

APIs (Application Programmer Interfaces) are the components that enable diverse platforms, apps, and systems to connect and share data with each other. Think of APIs as a set of software modules, tools, and protocols that enable two or more platforms, systems and most commonly, applications to communicate with each other and initiate tasks or processes. APIs are essential for defining and customizing Graphical User Interfaces (GUIs) too. Cloud platform providers all have extensive APIs defined and work in close collaboration with development partners to fine-tune app performance using them. [Amazon Web Services](#), [Facebook](#), [Google](#), [Marketo](#), [Salesforce](#), [SAP Hybris](#), [Twitter](#) and thousands of other companies have APIs available. As of today, the [Programmable Web](#) lists 16,590 APIs in its database.

- A large part of developing microservices and cloud applications is determining
 - Which APIs do I want to publish to 3rd parties or unanticipated calling applications.
 - Which 3rd party/cloud APIs do I want to call instead of writing the equivalent function?
 - Only implement a microservice if:
 - You cannot find a web API that meets your needs.
 - Your implementation will be differentiated and core to your business model.
- This is why I keep harping on the fact that you should not assume your UI is the only thing calling your APIs.
- We will do one example in this course, SmartyStreets, and I can show one other, Stripe.

SmartyStreets

SmartyStreets Sequence



SmartyStreets Demo(s) – UI

localhost:3000/e6156/index.html#!/profile

Apps CourseWorks pow... Settings MGRS Coordinate... Center Harbor, N... Best Online Cours... An interactive, mo... Home Page | Colu... https://app.graph... ENGI 1006 Spring

Welcome to E6156 Home About Services Contact Logout My Profile/My Account

Enter and Address

520 w 120th

- 520 W 120th St, New York NY
- 520 W 120th Pl, Alsip IL
- 520 W 120th St, Alsip IL
- 520 W 120th Dr, Alton KS
- 520 W 120th St, Bedford OH
- 520 W 120th St, Brookpark OH
- 520 W 120th St, Beachwood OH
- 520 W 120th St, Brecksville OH
- 520 W 120th St, Bay Village OH
- 520 W 120th St S, Bluff City KS



et,
consectetur adipisicing elit. Quod
tenetur ex natus at dolorem enim!
Nesciunt pariatur voluptatem sunt
quam eaque, vel, non in id dolore
voluptates quos eligendi labore.

Card Three

Lorem ipsum dolor sit amet,
consectetur adipisicing elit. Rem
magni quas ex numquam, maxime
minus quam molestias corporis
quod, ea minima accusamus.

UI – HTML/JavaScript

```
<!-- Call to Action Well -->


<div class="card-body">
    <h2>Enter and Address</h2>
    <form>
      <br>
      <input type="text" id="newaddress" name="newaddress" placeholder="{{placeholder}}" style="...">
      <br>
    </form>
    <p></p>
    <div class="dropdown">
      <button class="btn btn-primary dropdown-toggle"
             type="button" id="dropdownMenuButton" data-toggle="dropdown" aria-haspopup="true" aria-expanded="false" style="...">
        Choose the type of address
      </button>
      <div class="dropdown-menu" aria-labelledby="dropdownMenuButton">
        <a ng-click="addressKind($index)" ng-repeat="k in addressKinds" class="dropdown-item">{{k}}</a>
      </div>
    </div>
  </div>
</div>


```

UI JavaScript

Instead of "Boo Yah"

- POST to profile endpoint
- Address info.
- Address type.

```
CustomerApp.controller("ProfileController", function($scope, $http, $location, $window) {  
  console.log("Profile controller loaded.")  
  
  let s3 = jQuery.LiveAddress({  
    key: "18981749384552786",  
    waitForStreet: true,  
    debug: true,  
    target: "US",  
    placeholder: "Enter address",  
    addresses: [{  
      freeform: '#newaddress'  
    }]  
  });  
  
  s3.on("AddressAccepted", function(event, data, previousHandler)  
  {  
    console.log("Boo Yah!")  
  });  
  
  $scope.placeholder = "enter an address and select a choice."  
  $scope.addressKinds = ['Home', 'Work', 'Other']  
  
  $scope.addressKind = function(idx) {  
    console.log("Address kind = " + $scope.addressKinds[idx]);  
  };  
});
```

SmartyStreets – Microservice Verification Call

GET <https://us-street.api.smartystreets.com/street-address?street=520+w 120th+ny+ny&auth-id=cf5a1a2a-4a20-433a-8a20-000000000000>

Params **Send** **Save**

Authorization Headers (1) Body Pre-request Script Tests **CC**

Type No Auth

Body Cookies Headers (4) Test Results Status: 200 OK Time: 70 m

Pretty Raw Preview **JSON** 

```
1 [ { 2   "input_index": 0, 3   "candidate_index": 0, 4   "delivery_line_1": "520 W 120th St", 5   "last_line": "New York NY 10027-6601", 6   "delivery_point_barcode": "100276601205", 7   "components": { 8     "primary_number": "520", 9     "street_predirection": "W", 10    "street_name": "120th", 11    "street_suffix": "St", 12    "city_name": "New York", 13    "state_abbreviation": "NY", 14    "zipcode": "10027", 15    "plus4_code": "6601", 16    "delivery_point": "20", 17    "delivery_point_check_digit": "5" 18  }, 19  "metadata": { 20    "record_type": "S", 21    "zip_type": "Standard", 22    "county_fips": "36061", 23    "county_name": "New York", 24    "state_abbr": "NY" 25  } }
```

Example Code

```
auth_id = "c6eb7[REDACTED]cc-445f305eaf32"
auth_token = "wY[REDACTED]YM"
base_url = "https://us-street.api.smartystreets.com/street-address"

def check_address(street_no, street_name, town, state):

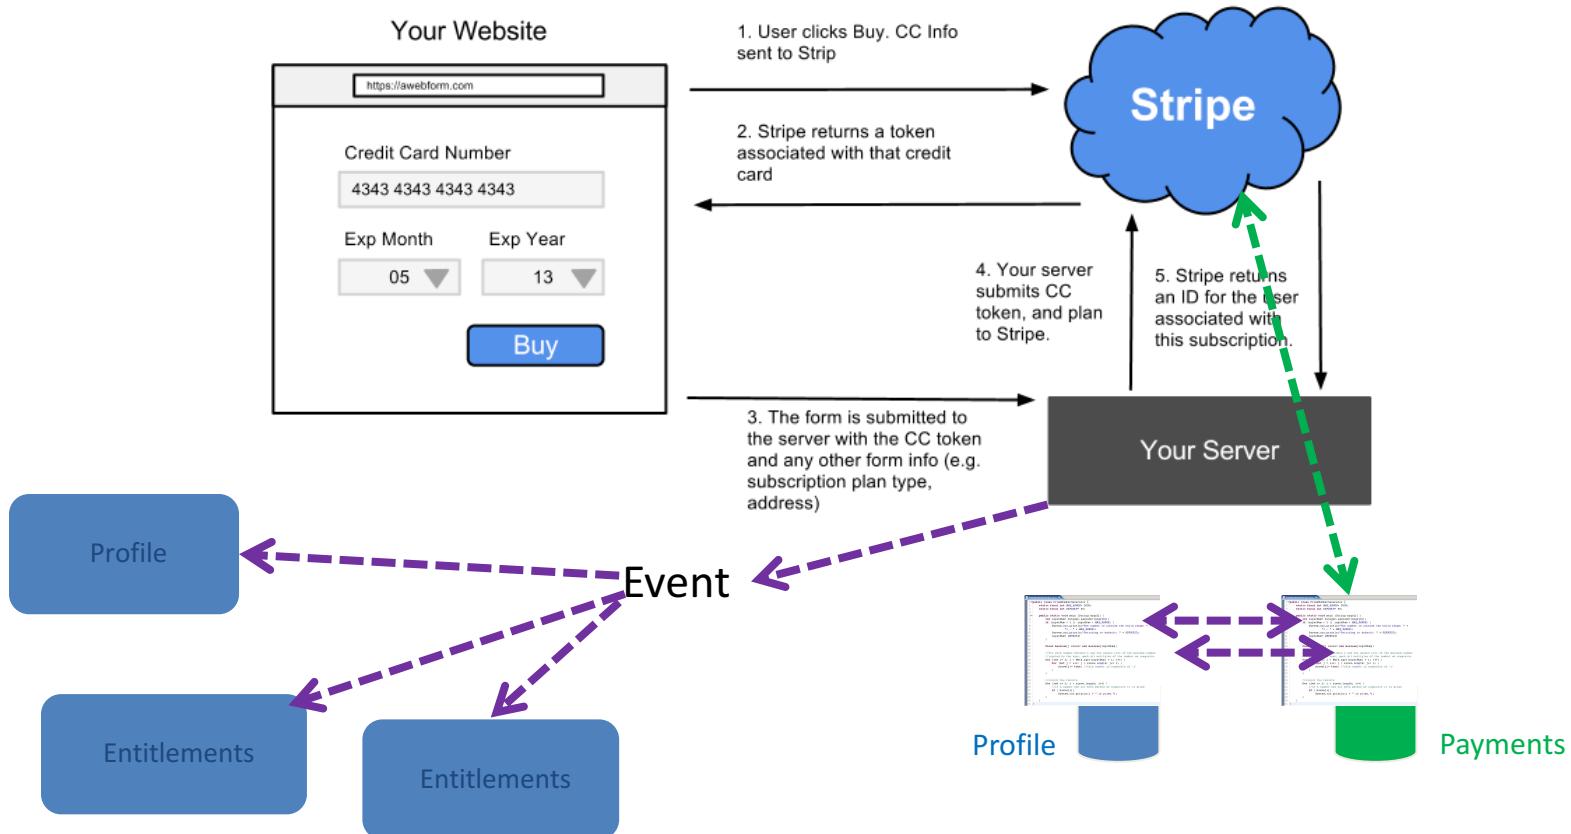
    get_url = base_url + "?street="
    added = False
    if street_no:
        added = True
        get_url += str(street_no)
    if street_name:
        if added:
            get_url += "+"
        get_url += street_name
    if town:
        if added:
            get_url += "+"
        get_url += town
    if state:
        if added:
            get_url += "+"
        get_url += state

    if added:
        get_url += "&auth-id=" + auth_id + "&auth-token=" + auth_token

    print("URL = ", get_url)
    r = requests.get(get_url)
    return r
```

Stripe

Stripe



Publish/Subscribe

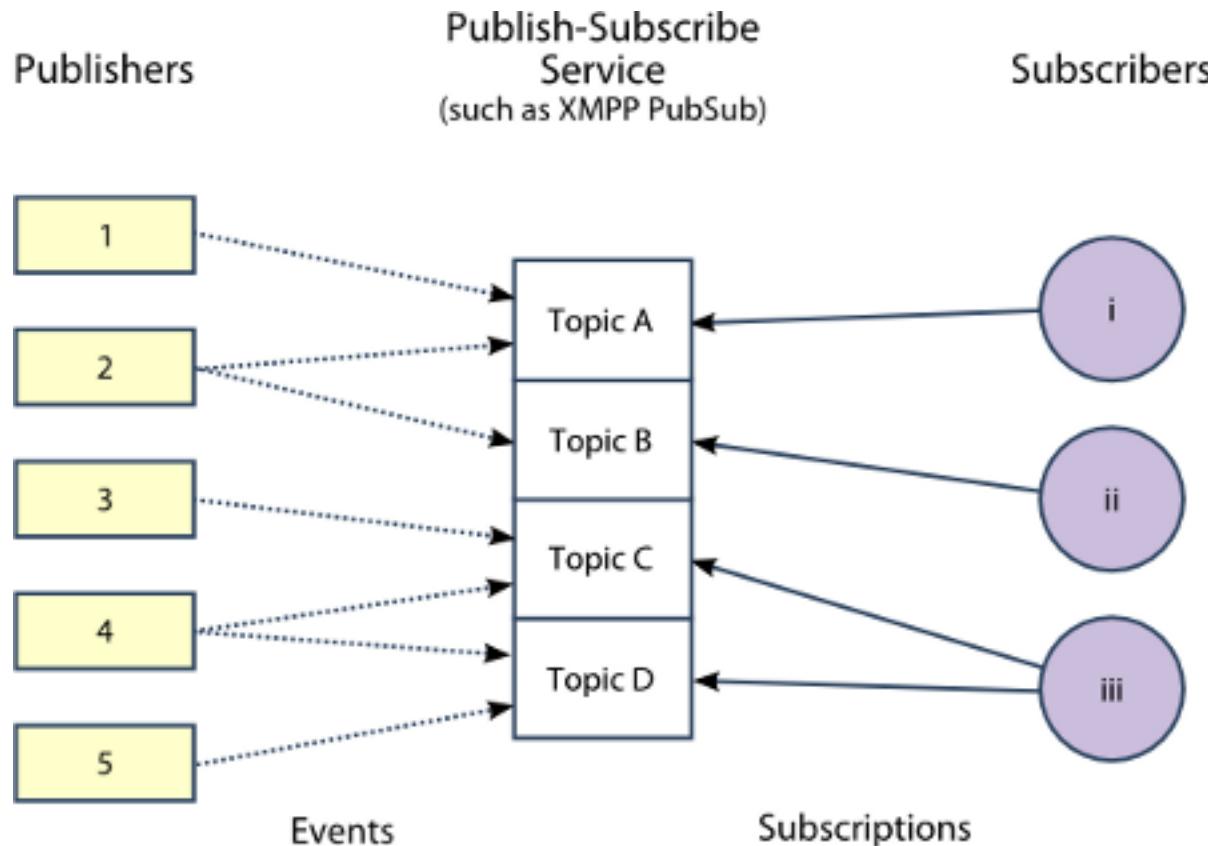
Publish/Subscribe (Wikipedia)

“In [software architecture](#), **publish–subscribe** is a [messaging pattern](#) where senders of [messages](#), called publishers, do not program the messages to be sent directly to specific receivers, called subscribers, but instead categorize published messages into classes without knowledge of which subscribers, if any, there may be. Similarly, subscribers express interest in one or more classes and only receive messages that are of interest, without knowledge of which publishers, if any, there are.

Publish–subscribe is a sibling of the [message queue](#) paradigm, and is typically one part of a larger [message-oriented middleware](#) system. Most messaging systems support both the pub/sub and message queue models in their [API](#), e.g. [Java Message Service](#) (JMS).

This pattern provides greater network [scalability](#) and a more dynamic [network topology](#), with a resulting decreased flexibility to modify the publisher and the structure of the published data.”

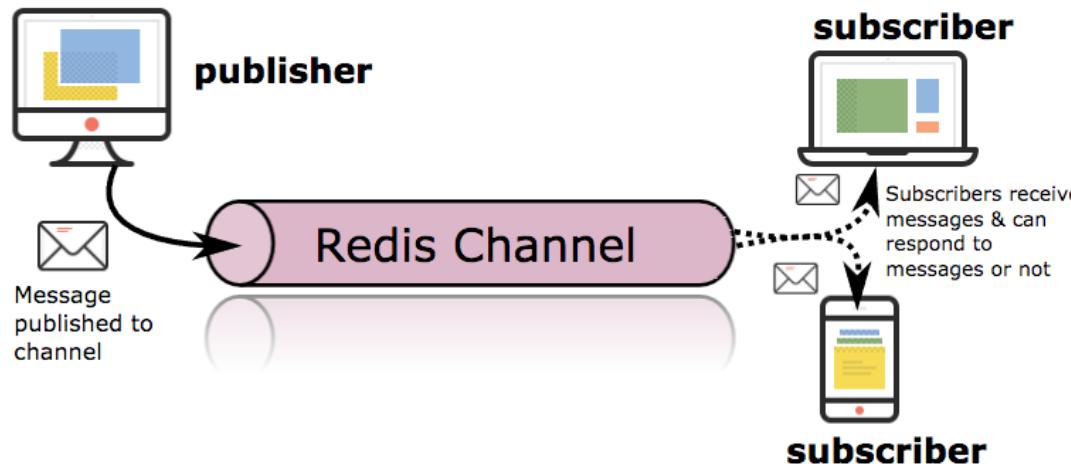
Publish/Subscribe



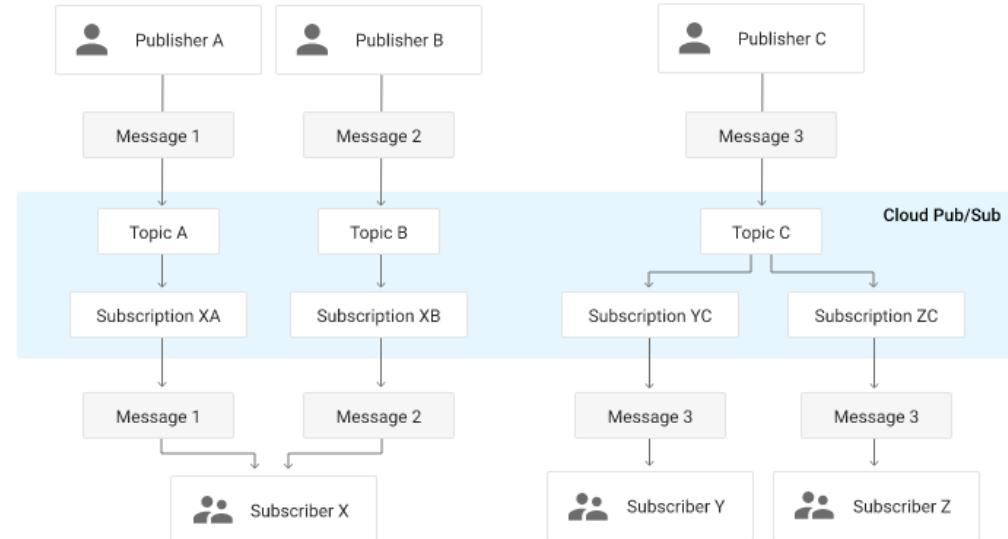
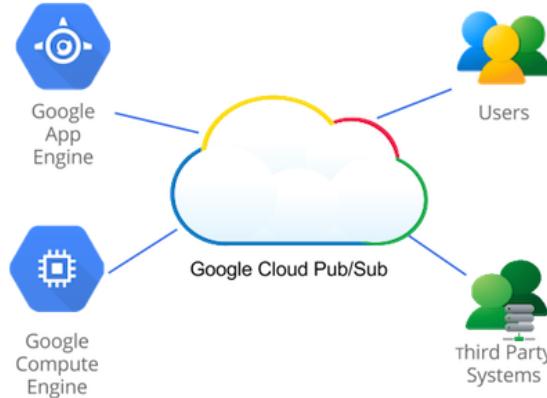
Redis Example

Pub/Sub Messaging Pattern

Redis Pub/Sub uses a message passing system that message senders - called **publishers** - post a message to a channel that the message receivers - called **subscribers** - can respond to messages without either the publishers or subscribers knowing any details about each other.

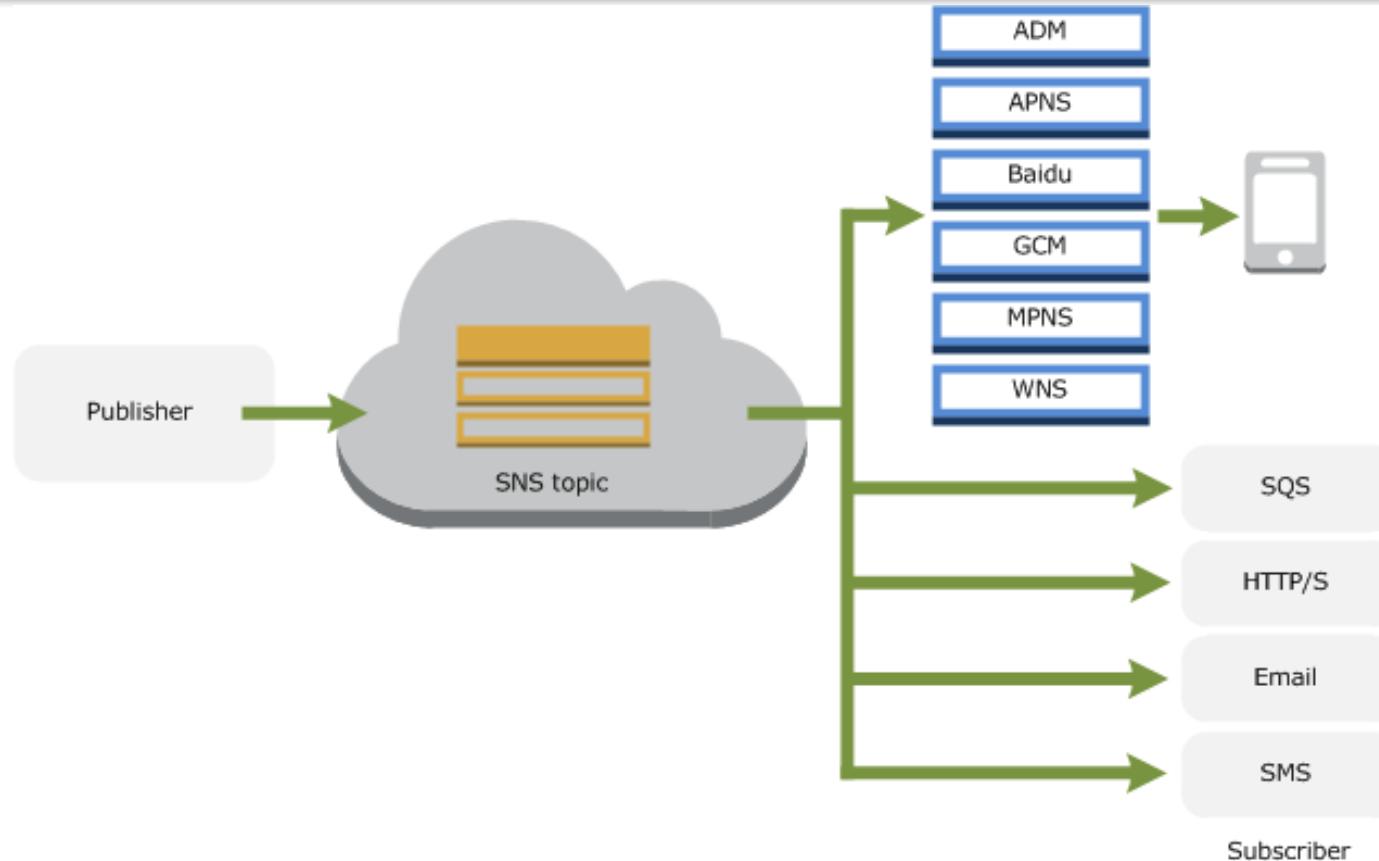


Google Cloud Pub/Sub



<https://cloud.google.com/pubsub/overview>

SNS



Publish/Subscribe

Message Queueing (Wikipedia)

"Message queues provide an [asynchronous communications protocol](#), meaning that the sender and receiver of the message do not need to interact with the message queue at the same time. Messages placed onto the queue are stored until the recipient retrieves them. Message queues have implicit or explicit limits on the size of data that may be transmitted in a single message and the number of messages that may remain outstanding on the queue.

... ...

There are often numerous options as to the exact semantics of message passing, including:

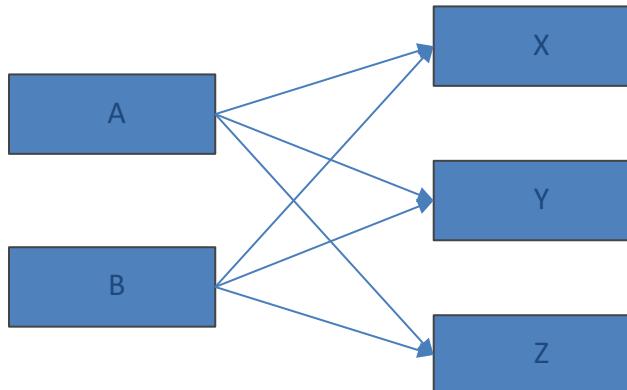
- Durability - messages may be kept in memory, written to disk, or even committed to a [DBMS](#) if the need for reliability indicates a more resource-intensive solution.
- Security policies - which applications should have access to these messages?
- Message purging policies - queues or messages may have a "[time to live](#)"
- Message filtering - some systems support filtering data so that a subscriber may only see messages matching some pre-specified criteria of interest
- Delivery policies - do we need to guarantee that a message is delivered at least once, or no more than once?
- Routing policies - in a system with many queue servers, what servers should receive a message or a queue's messages?
- Batching policies - should messages be delivered immediately? Or should the system wait a bit and try to deliver many messages at once?
- Queuing criteria - when should a message be considered "enqueued"? When one queue has it? Or when it has been forwarded to at least one remote queue? Or to all queues?
- Receipt notification - A publisher may need to know when some or all subscribers have received a message.

Message Queue Service Pattern

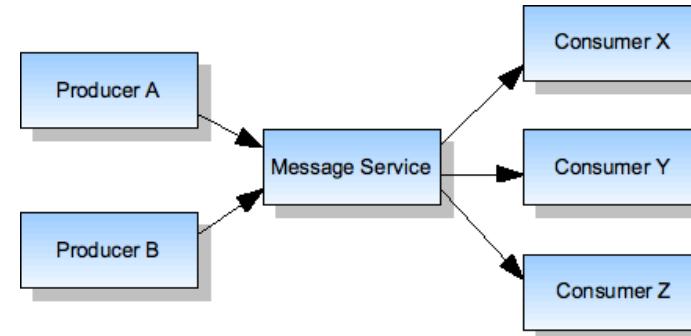
Point-to-Point communication between modules becomes fragile at scale (no. of messages, no. of participants)

- Adding modules P and Q requires finding and configuring all senders.
- I want to send M to any one of X1, X2 and X3, but only one.
- I want to make sure that someone processes M but do not want to hold the transaction until I get a response (I may not even need the response).
- My destinations are not “up” at the same times I am.

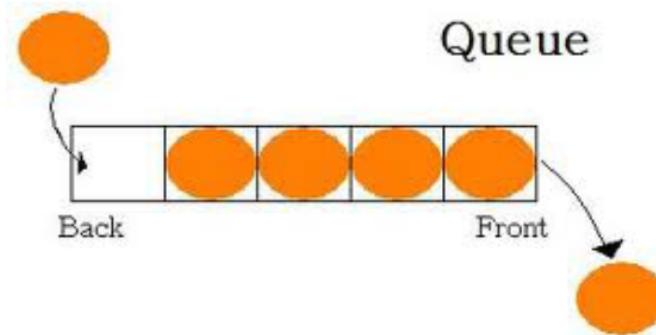
Anti-Pattern



Best Practice Pattern



Queue?



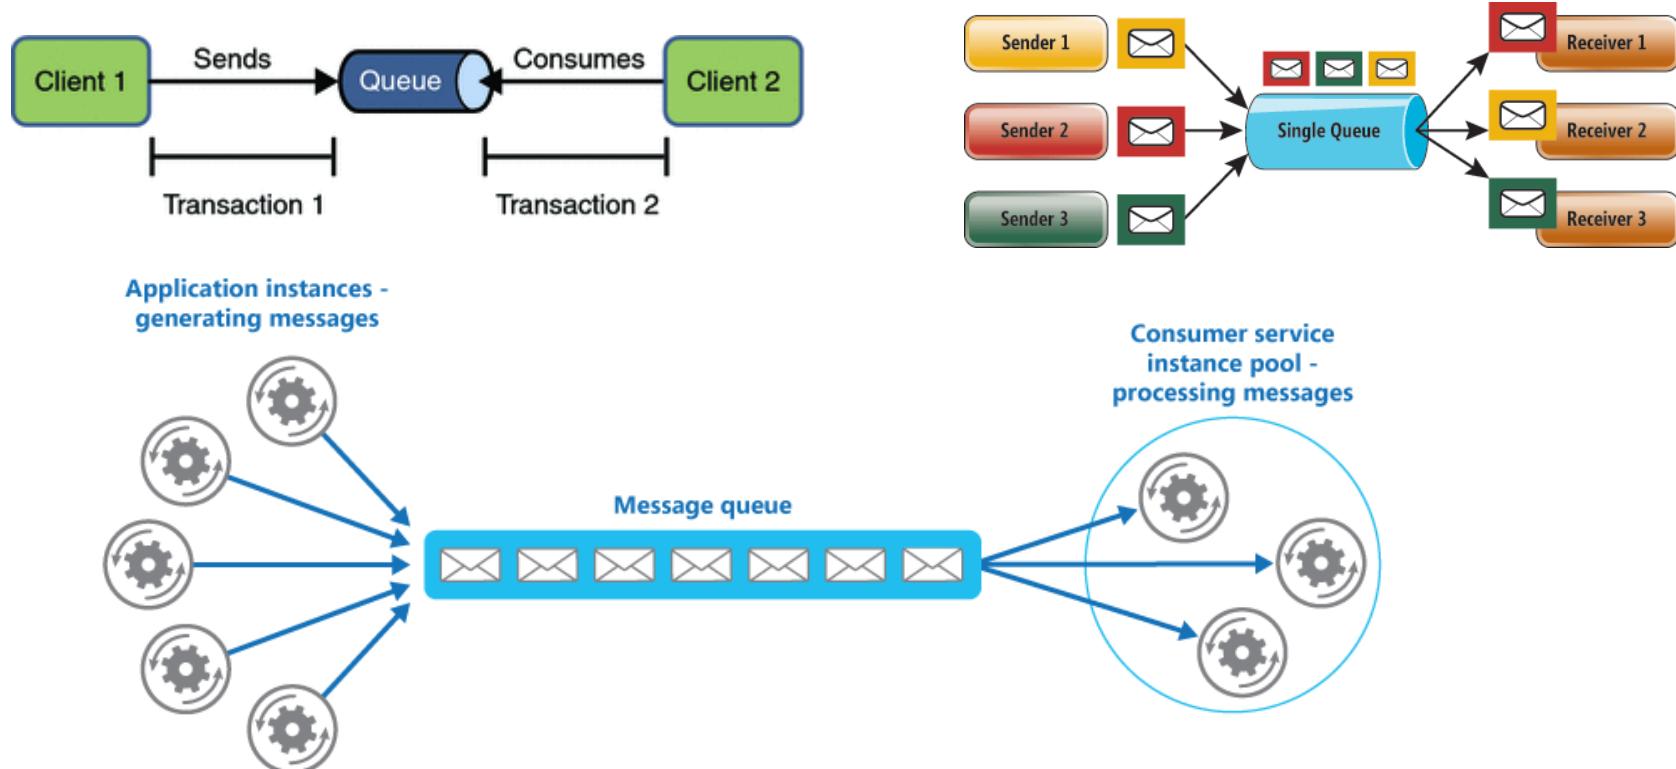
Queue Operations

- Create Q
- Delete Q
- Send Message
- Receive Message
- Delete Message

FIFO Semantics

- Strict, limited scalability.
- Best effort, highly scalable.
- Transactional

Message Exchange Patterns



Why Would I Use a Queue?

<http://www.javacodegeeks.com/2013/06/working-with-amazon-simple-queue-service-using-java.html>

- Flow control
 - My server may be able to process 100 reqs/second
 - There are thousands of clients
 - The aggregate rate could be 10 req/s or 300 req/s
 - Without a queue, clients would get connect failures under load
- The request is going to take a long time
 - Image reformatting and tagging
 - I cannot have my code “wait” for a response
 - And I do not want to write a lot of “Is it done yet?” calls
- I do not know (or care) which one implements the logic
 - I have 5 – 10 image processors for GIFs
 - I have 7 – 10 processors for .wav files
 - I do not want to know which ones are active, when and processing what.
 - I just want to put the “thing in a queue” knowing that someone will eventually pick it up.
- Think “email” for API calls

Amazon SQS – Example

